

Linux Device Driver Tutorial Part 27 – Using High Resolution Timer In Linux Device Driver

Post Contents [[hide](#)]

1 High Resolution Timer (HRT/hrtimer)

2 Users of High Resolution Timer

3 High Resolution timer API

3.1 ktime_set

3.2 Initialize High Resolution Timer

3.2.1 hrtimer_init

3.3 Start High Resolution Timer

3.3.1 hrtimer_start

3.4 Stop High Resolution Timer

3.4.1 hrtimer_cancel

3.4.2 hrtimer_try_to_cancel

3.5 Changing the High Resolution Timer's Timeout

3.5.1 hrtimer_forward

3.5.2 hrtimer_forward_now

3.6 Check High Resolution Timer's status

3.6.1 hrtimer_get_remaining

3.6.2 hrtimer_callback_running

3.6.3 hrtimer_cb_get_time

4 Using High Resolution Timer In Linux Device Driver

4.1 Driver Source Code

5 Building and Testing Driver

6 Points to remember

6.0.1 Share this:

6.0.2 Like this:

6.0.3 Related

High Resolution Timer (HRT/hrtimer)

In our [last tutorial](#) we have seen kernel timer. Now we are taking about high resolution timer. Everyone might have some questions. Why the hell we need two timers? Why can they merge two timers into one? Can't able to integrate? Yes. They have tried to merge these two timers. But they have failed. Because **Cascading Timer Wheel (CTW)** is used in kernel timer. Cascading Timer Wheel (CTW) code is fundamentally not suitable for such an approach like merging these two timers. Because hrtimer is maintaining a time-ordered data

structure of timers (timers are inserted in time order to minimize processing at activation time). The data structure used is a red-black tree, which is ideal for performance-focused applications (and happens to be available generically as a library within the kernel).

Kernel Timers are bound to **jiffies**. But this High Resolution Timer (HRT) is bound with 64-bit **nanoseconds** resolution.

With kernel version 2.6.21 onwards, high resolution timers (HRT) are available under Linux. For this, the kernel has to be compiled with the configuration parameter **CONFIG_HIGH_RES_TIMERS** enabled.

There are many ways to check whether high resolution timers are available,

- In the **/boot** directory, check the kernel config file. It should have a line like **CONFIG_HIGH_RES_TIMERS=y**.
- Check the contents of **/proc/timer_list**. For example, the **.resolution** entry showing 1 nanosecond and event_handler as **hrtimer_interrupt** in **/proc/timer_list** indicate that high resolution timers are available.
- Get the clock resolution using the **clock_getres** system call.

Users of High Resolution Timer

- The primary users of precision timers are user-space applications that utilize **nanosleep**, **posix-timers** and **Interval Timer (itimer)** interfaces.
- In-kernel users like drivers and subsystems which require precise timed events (e.g. multimedia).

High Resolution timer API

We need to include the **<linux/hrtimer.h>** (**#include <linux/hrtimer.h>**) in order to use kernel timers. Kernel timers are described by the **hrtimer** structure, defined in **<linux/hrtimer.h>**:

```
struct hrtimer {  
    struct rb_node node;  
    ktime_t expires;  
    int (* function) (struct hrtimer *);  
    struct hrtimer_base * base;  
};
```

Where,

node – red black tree node for time ordered insertion

expires – the absolute expiry time in the hrtimers internal representation. The time is related to the clock on which the timer is based.

function – timer expiry callback function. This function has an integer return value, which should be either HRTIMER_NORESTART (for a one-shot timer which should not be started again) or HRTIMER_RESTART for a recurring timer. In the restart case, the callback must set a new expiration time before returning.

base – pointer to the timer base (per cpu and per clock)

The **hrtimer** structure must be initialized by `init_hrtimer_#CLOCKTYPE`.

ctime_set

There is a new type, **ctime_t**, which is used to store a time value in nanoseconds. On 64-bit systems, a **ctime_t** is really just a 64-bit integer value in nanoseconds. On 32-bit machines, however, it is a two-field structure: one 32-bit value holds the number of seconds, and the other holds nanoseconds. The below function is used to get the **ctime_t** from seconds and nanoseconds.

ctime_set(long secs, long nanosecs);

Arguments:

secs – seconds to set

nsecs – nanoseconds to set

Return:

The **ctime_t** representation of the value.

Initialize High Resolution Timer

hrtimer_init

void hrtimer_init(struct hrtimer *timer, clockid_t clock_id, enum hrtimer_mode mode);

Arguments:

timer – the timer to be initialized

clock_id – the clock to be used

The clock to use is defined in `./include/linux/time.h` and represents the various clocks that the system supports (such as the real-time clock or a monotonic clock that simply represents time from a starting point, such as system boot).

CLOCK_MONOTONIC: a clock which is guaranteed always to move forward in time, but which does not reflect “wall clock time” in any specific way. In the

current implementation, *CLOCK_MONOTONIC* resembles the jiffies tick count in that it starts at zero when the system boots and increases monotonically from there.

CLOCK_REALTIME: which matches the current real-world time.

mode – timer mode absolute (HRTIMER_MODE_ABS) or relative (HRTIMER_MODE_REL)

Start High Resolution Timer

Once a timer has been initialized, it can be started with the below mentioned function.

hrtimer_start

```
inthrtimer_start(struct hrtimer *timer, ktime_t time, const enum hrtimer_mode mode);
```

This call is used to (Re)start an hrtimer on the current CPU.

Arguments:

timer – the timer to be added

time – expiry time

mode – expiry mode: absolute (HRTIMER_MODE_ABS) or relative (HRTIMER_MODE_REL)

Returns:

0 on success 1 when the timer was active

Stop High Resolution Timer

Using below function, we can able to stop the High Resolution Timer.

hrtimer_cancel

```
inthrtimer_cancel(struct hrtimer * timer);
```

This will cancel a timer and wait for the handler to finish.

Arguments:

timer – the timer to be cancelled

Returns:

- 0 when the timer was not active
- 1 when the timer was active

hrtimer_try_to_cancel

int hrtimer_try_to_cancel (struct hrtimer * timer);

This will try to deactivate a timer.

Arguments:

timer – hrtimer to stop

Returns:

- 0 when the timer was not active
- 1 when the timer was active
- -1 when the timer is currently executing the callback function and cannot be stopped

Changing the High Resolution Timer's Timeout

If we are using this High Resolution Timer (hrtimer) as periodic timer, then the callback must set a new expiration time before returning. Usually, restarting timers are used by kernel subsystems which need a callback at a regular interval.

hrtimer_forward

u64 hrtimer_forward (struct hrtimer * timer, ktime_t now, ktime_t interval);

This will forward the timer expiry so it will expire in the future by the given interval.

Arguments:

timer – hrtimer to forward

now – forward past this time

interval – the interval to forward

Returns:

Returns the number of overruns.

hrtimer_forward_now

u64 hrtimer_forward_now(struct hrtimer *timer, ktime_t interval);

This will forward the timer expiry so it will expire in the future from now by the given interval.

Arguments:

timer – hrtimer to forward

interval – the interval to forward

Returns:

Returns the number of overruns.

Check High Resolution Timer's status

The below explained functions are used to get the status and timings.

hrtimer_get_remaining

ktime_thrtimer_get_remaining(const struct hrtimer * timer);

This is used to get remaining time for the timer.

Arguments:

timer – hrtimer to get the remaining time

Returns:

Returns the remaining time.

hrtimer_callback_running

int hrtimer_callback_running(struct hrtimer *timer);

This is the helper function to check, whether the timer is running the callback function.

Arguments:

timer – hrtimer to check

Returns:

- 0 when the timer's callback function is not running
- 1 when the timer's callback function is running

hrtimer_cb_get_time

ktime_thrtimer_cb_get_time(struct hrtimer *timer);

This function used to get the current time of the given timer.

Arguments:

timer – hrtimer to get the time

Returns:

Returns the time.

Using High Resolution Timer In Linux Device Driver

In this example we took the basic driver source code from [this](#) tutorial. On top of that code we have added the high resolution timer. The steps are mentioned below.

1. Initialize and start the timer in init function
2. After timeout, registered timer callback will be called.
3. In the timer callback function again we are forwarding the time period and return **HRTIMER_RESTART**. We have to do this step if we want periodic timer. Otherwise we can ignore that time forwarding and return **HRTIMER_NORESTART**.
4. Once we are done, we can disable the timer.

Driver Source Code

driver.c:

```
#include <linux/kernel.h>
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kdev_t.h>
#include <linux/fs.h>
#include <linux/cdev.h>
#include <linux/device.h>
#include <linux/hrtimer.h>
#include <linux/ktime.h>

//Timer Variable
#define TIMEOUT 5000 * 1000000L //nano seconds
static struct hrtimer tx_hr_timer;
static unsigned int count = 0;

dev_t dev = 0;
static struct class *dev_class;
static struct cdev tx_cdev;

static int __init tx_driver_init(void);
```

```
static void __exit etx_driver_exit(void);
static int etx_open(struct inode *inode, struct file *file);
static int etx_release(struct inode *inode, struct file *file);
static ssize_t etx_read(struct file *filp, char __user *buf, size_t len, loff_t * off);
static ssize_t etx_write(struct file *filp, const char *buf, size_t len, loff_t * off);
```

```
static struct file_operations fops =
{
    .owner      = THIS_MODULE,
    .read       = etx_read,
    .write      = etx_write,
    .open       = etx_open,
    .release    = etx_release,
};
```

```
//Timer Callback function. This will be called when timer expires
enum hrtimer_restart timer_callback(struct hrtimer *timer)
{
    /* do your timer stuff here */
    printk(KERN_INFO "Timer Callback function Called [%d]\n", count++);
    hrtimer_forward_now(timer, ktime_set(0, TIMEOUT));
    return HRTIMER_RESTART;
}
```

```
static int etx_open(struct inode *inode, struct file *file)
{
    printk(KERN_INFO "Device File Opened...!!!\n");
    return 0;
}
```

```
static int etx_release(struct inode *inode, struct file *file)
{
    printk(KERN_INFO "Device File Closed...!!!\n");
    return 0;
}
```

```
static ssize_t etx_read(struct file *filp, char __user *buf, size_t len, loff_t * off)
```



```

{
printk(KERN_INFO "Read Function\n");
return 0;
}
static ssize_t etx_write(struct file *filp, const char __user *buf, size_t len, loff_t
*off)
{
printk(KERN_INFO "Write function\n");
return 0;
}

static int __init etx_driver_init(void)
{
ktime_t ktime;

    /*Allocating Major number*/
    if((alloc_chrdev_region(&dev, 0, 1, "etx_Dev")) < 0){
printk(KERN_INFO "Cannot allocate major number\n");
return -1;
    }
    printk(KERN_INFO "Major = %d Minor = %d \n", MAJOR(dev), MINOR(dev));

    /*Creating cdev structure*/
    cdev_init(&etx_cdev, &fops);

    /*Adding character device to the system*/
    if((cdev_add(&etx_cdev, dev, 1)) < 0){
printk(KERN_INFO "Cannot add the device to the system\n");
goto _class;
    }

    /*Creating struct class*/
    if((dev_class = class_create(THIS_MODULE, "etx_class")) == NULL){
printk(KERN_INFO "Cannot create the struct class\n");
goto _class;
    }
}

```

```

    /*Creating device*/
    if((device_create(dev_class,NULL,dev,NULL,"etx_device")) == NULL){
        printk(KERN_INFO "Cannot create the Device 1\n");
        goto _device;
    }

    ktime = ktime_set(0, TIMEOUT);
    hrtimer_init(&etx_hr_timer, CLOCK_MONOTONIC, HRTIMER_MODE_REL);
    etx_hr_timer.function = &timer_callback;
    hrtimer_start(&etx_hr_timer, ktime, HRTIMER_MODE_REL);

    printk(KERN_INFO "Device Driver Insert...Done!!!\n");
    return 0;
_r_device:
    class_destroy(dev_class);
_r_class:
    unregister_chrdev_region(dev,1);
    return -1;
}

void __exit etx_driver_exit(void)
{
    //stop the timer
    hrtimer_cancel(&etx_hr_timer);
    device_destroy(dev_class,dev);
    class_destroy(dev_class);
    cdev_del(&etx_cdev);
    unregister_chrdev_region(dev, 1);
    printk(KERN_INFO "Device Driver Remove...Done!!!\n");
}

module_init(etx_driver_init);
module_exit(etx_driver_exit);

MODULE_LICENSE("GPL");
MODULE_AUTHOR("EmbeTronicX<embetronicx@gmail.com>");
MODULE_DESCRIPTION("A simple device driver - High Resolution Timer");

```

```
MODULE_VERSION("1.22");
```

Makefile:

```
obj-m += driver.o
```

```
KDIR = /lib/modules/$(shell uname -r)/build
```

```
all:
```

```
make -C $(KDIR) M=$(shell pwd) modules
```

```
clean:
```

```
make -C $(KDIR) M=$(shell pwd) clean
```

Building and Testing Driver

- Build the driver by using Makefile (**sudo make**)
- Load the driver using **sudo insmod driver.ko**
- Now see the Dmesg (**dmesg**)

```
linux@embetronicx-VirtualBox: dmesg
```

```
[          2643.773119]          Device      Driver      Insert...Done!!!
[    2648.773546]      Timer      Callback      function      Called      [0]
[    2653.773609]      Timer      Callback      function      Called      [1]
[    2658.774170]      Timer      Callback      function      Called      [2]
[    2663.773271]      Timer      Callback      function      Called      [3]
[ 2668.773388] Timer Callback function Called [4]
```

- See the timestamp. That callback function is executing every 5 seconds.
- Unload the module using **sudo rmmod driver**

Points to remember

This timer callback function will be executed from interrupt context. If you want to check that, you can use function **in_interrupt()**, which takes no parameters and returns nonzero if the processor is currently running in interrupt context, either hardware interrupt or software interrupt. Since it is running in interrupt context, user cannot perform some actions inside the callback function mentioned below.

- Go to sleep or relinquish the processor
- Acquire a mutex
- Perform time-consuming tasks
- Access user space virtual memory